

Delaunay Triangulation in Parallel

Adarsh Prakash


CSE 633 : Parallel Algorithms

(Spring 2017)


Instructor: Dr. Russ Miller

Definition

Triangle ΔABC is a **Delaunay Triangle**, if no other points lie in the circumcircle formed by ΔABC .




Invalid Delaunay Δ



Valid Delaunay Δ s

Delaunay – Voronoi: Duality



A **Voronoi diagram** is constructed by connecting centers of all the circumcircles formed by the Delaunay Triangles in a graph.



Direct Applications

- Nearest Neighbor
- Graph Locality / Point Location
- Surface Mapping / Reconstruction
- Game Development
- Motion Capture
- Path Planning (Autonomous Navigation)
- Physics – studying forces..
- Chemistry – atomic charges..
- Biology, Astrophysics and so on.


Applications




Algorithm

- **Divide-and-conquer** algorithm proposed by **Leonidas Guibas** and **Jorge Stolfi** [1].
- Follows closely the Voronoi construction algorithm from **Shamos** and **Hoey** [2].
- Difference is it clearly describes how to make use of quad-edge data structure to avoid computation of complete hull.
- Properties:
 - A quad-edge knows its direction (origin-destination NOT point-point)
 - A quad-edge maintains pointers to all edges leaving from and terminating at their origin and destination. (4-8 pointers depending on implementation)
- Objective is to parallelize this algorithm.


Algorithm: Merge Step



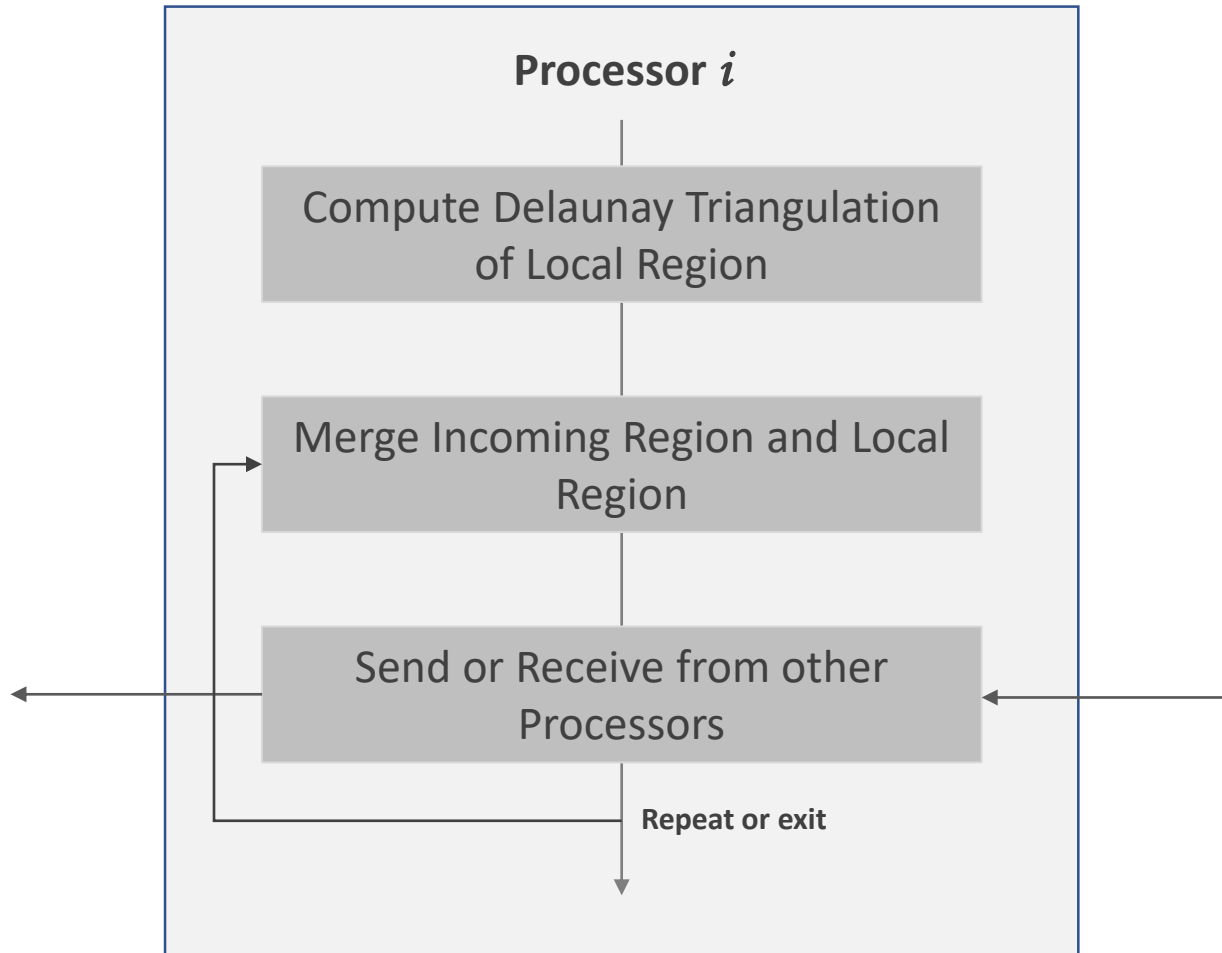
Algorithm: Merge Step




Algorithm: Merge Step



Algorithm: Parallel Overview



Domain Decomposition



Input space divided equally by X-Coordinate among Processors

Implementation

- Implementation in C and MPI
- Pseudo code from paper for serial version of merge – made life easier
- Jobs were run on *general-compute* and *largemem* partitions of CCR
- All communications are point-to-point: **MPI_Send** and **MPI_Recv**
- Data send/receive happens in a single block (as many as 31 million edges ~ 700mb)
- Approx. 500 jobs to cluster

Implementation

- **Input:**
 - Randomly generated points – **Bivariate Uniform Distribution** using Python **numpy** package
 - Equal range and density across both the axes
 - No duplicates and **pre-sorted** by **X-Coordinate**
 - Each coordinate is “**double**” precision $\in [0, 200*n]$
- **Output:**
 - Edge endpoints as indices

```
0, 162.422299106, 626335123.072
1, 235.609542392, 21674347.1286
2, 348.128895741, 545885503.786
3, 388.434040826, 160544722.935
...
```

Sample Input

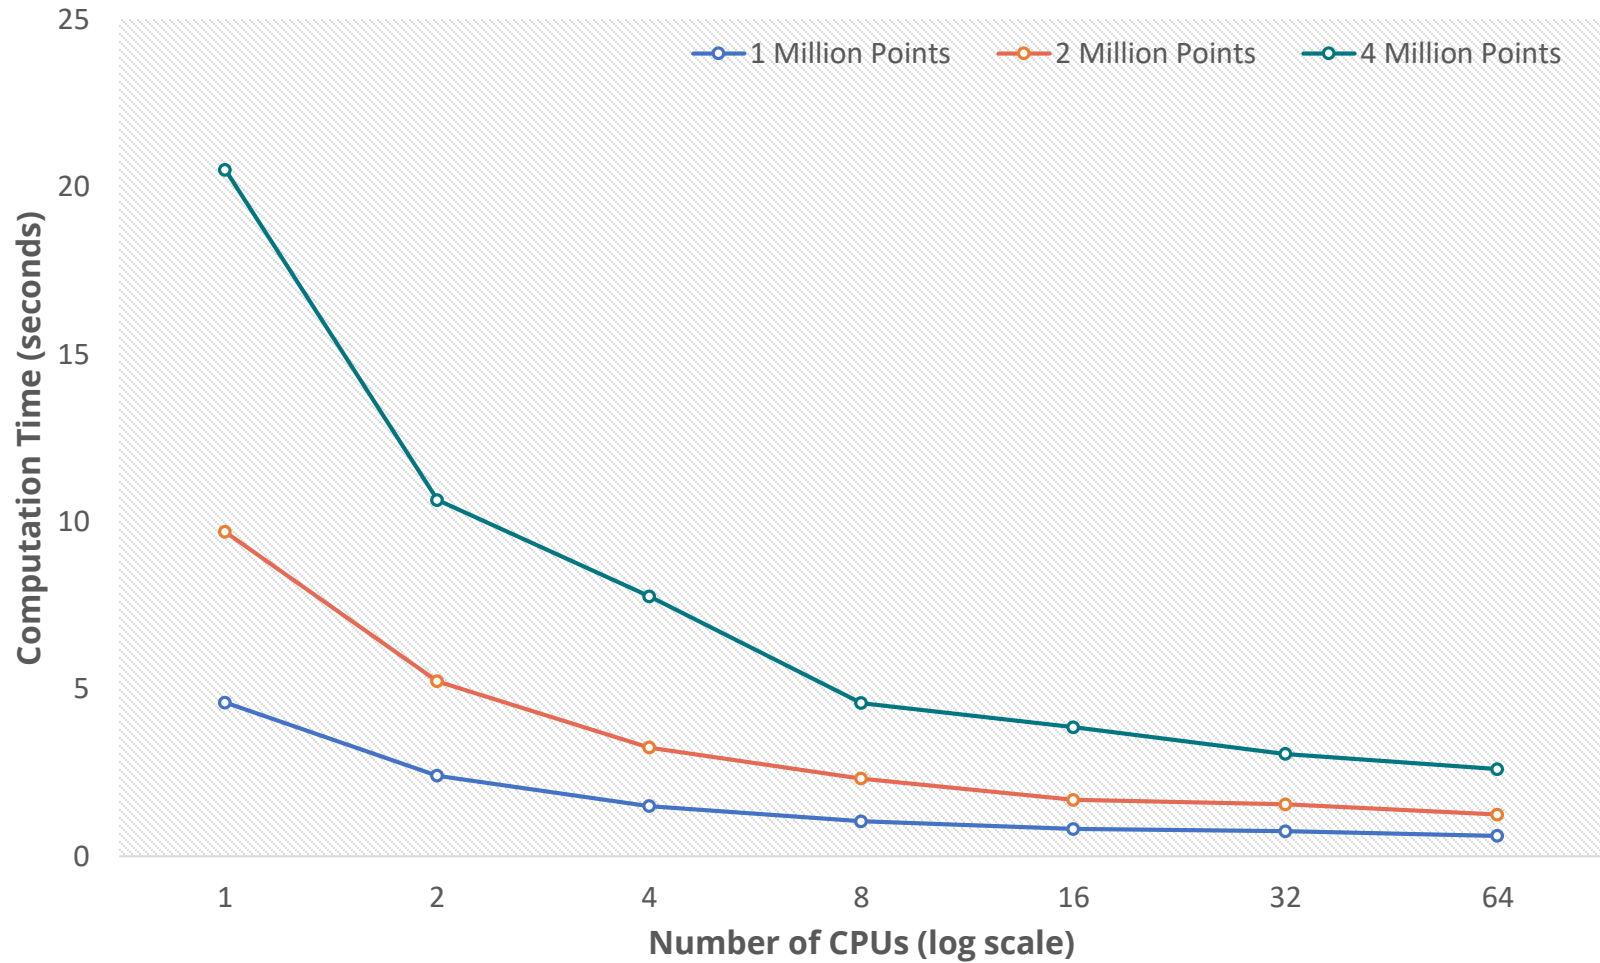
```
0 1
0 3
0 4
0 2
...
```

Sample Output

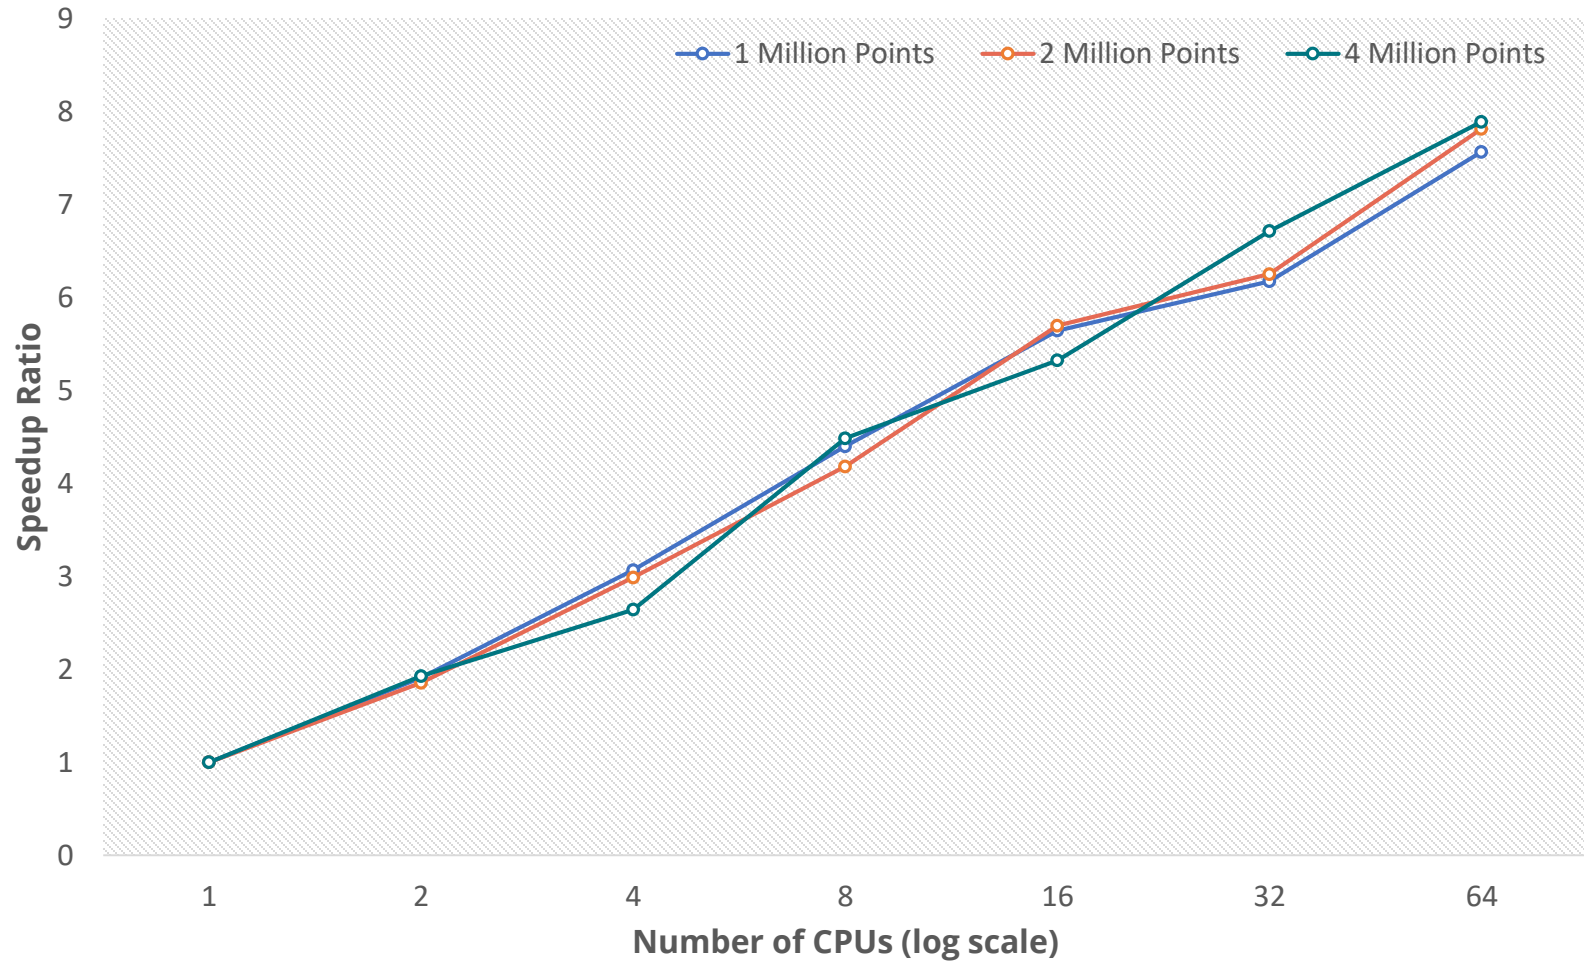
Results

- Run-times averaged over 3-5 jobs/runs
- Tried for several core-node combinations:
 - 2 CPUs per node with **shm** (intranode) and **tmi** (internode)
 - 1 CPU per node with **dapl** (internode)
 - 1 CPU per node with **tmi** (internode)
 - Upto 32 CPUs per node with **tcp** (intranode) and **tcp** (internode) – **I_MPI_FABRICS** and **I_MPI_FALLBACK** to the rescue!
- All results **validated** against results from standard packages:
 - Python (scipy.spatial.Delaunay) - faster
 - Matlab (triangulation)

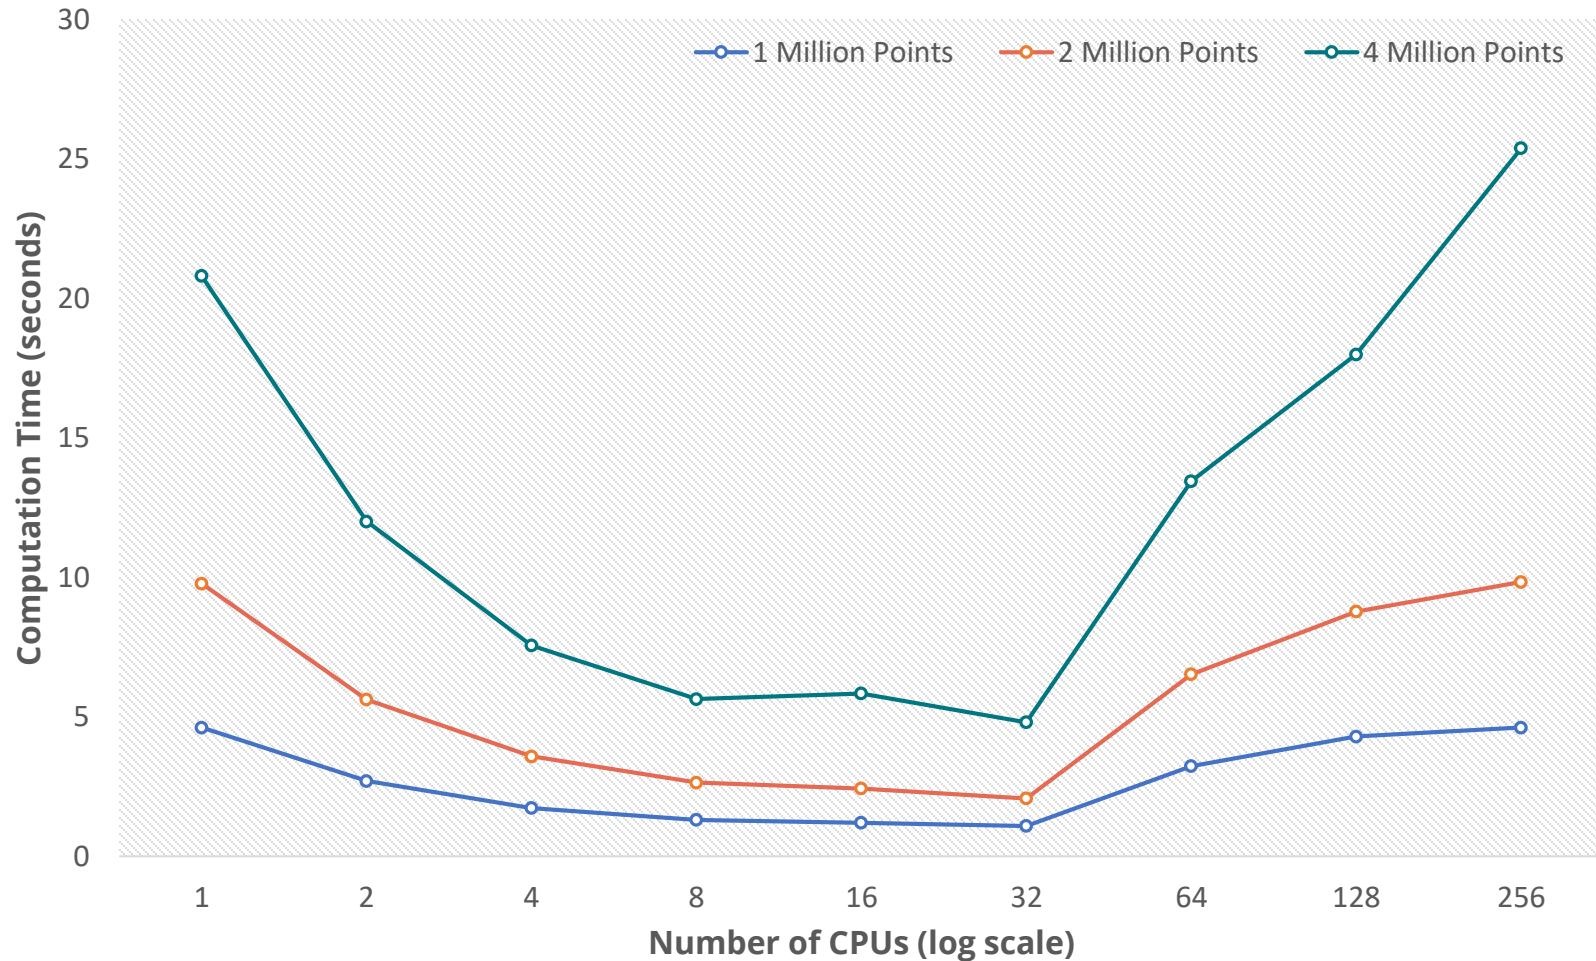
Time v/s CPU (1 CPU per node – TMI)



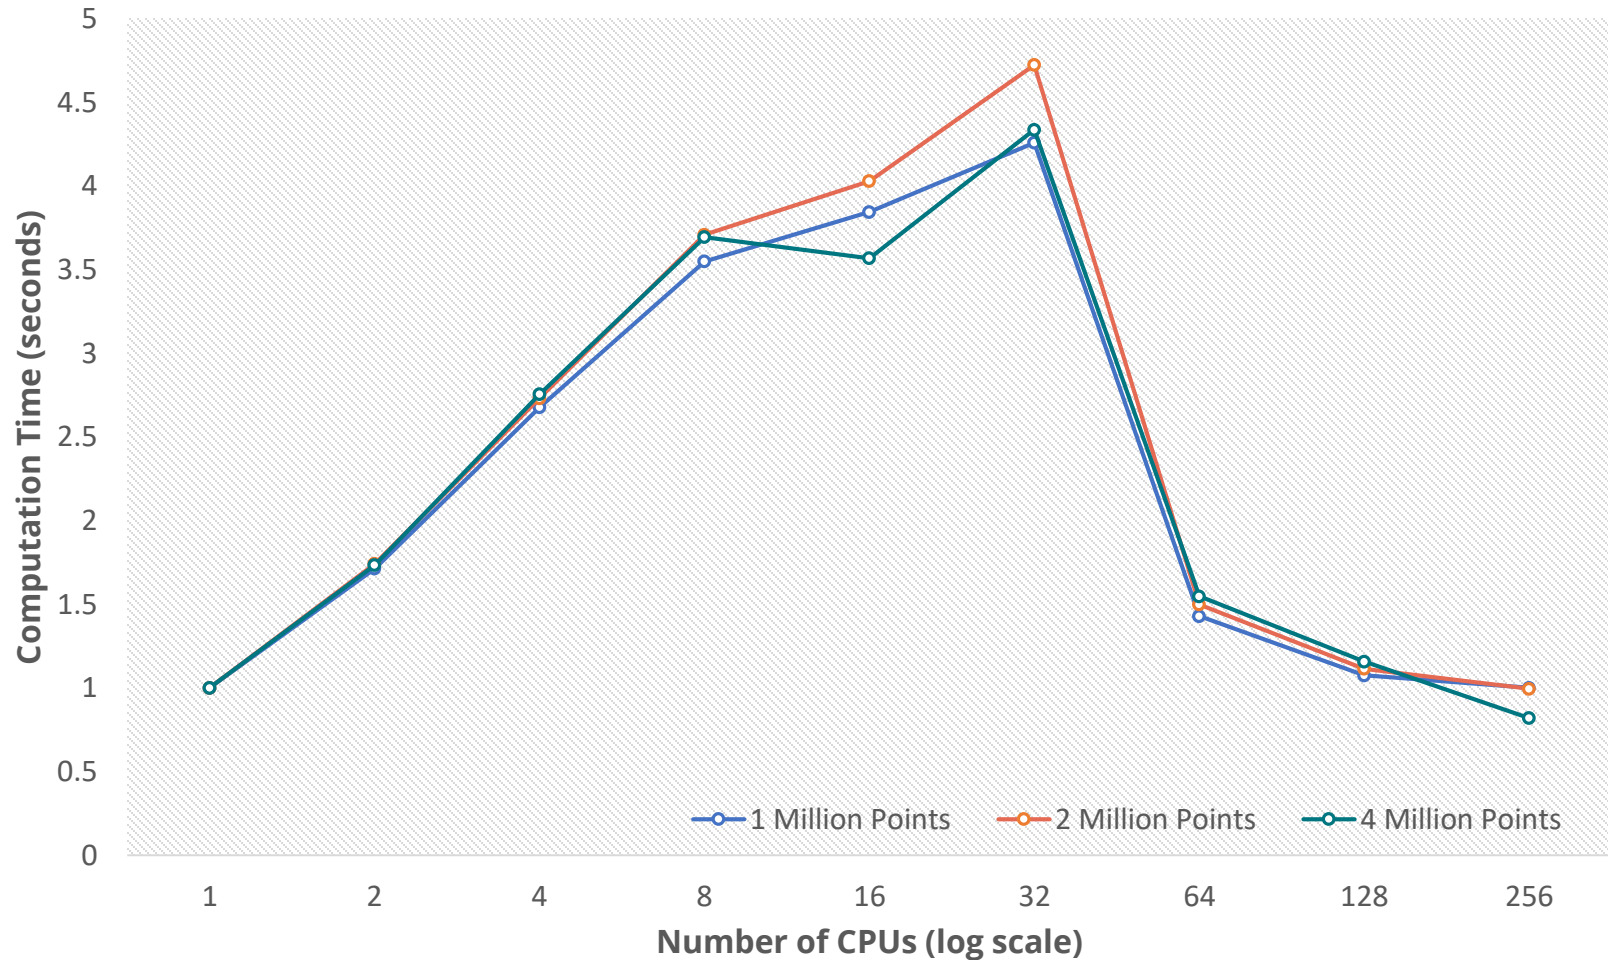
Speedup v/s CPU (1 CPU per node – TMI)




Time v/s CPU (32 CPUs per node – TCP – no shm)



Speedup v/s CPU (32 CPUs per node – TCP – no shm)



Asymptotic Growth (8 CPUs with 1 CPU per node)



Conclusion

- That drop in speedup for 32-cpus-per-node?
 - **Communication Cost: Intranode < Internode**
 - Difference is significant for TCP and hence the sudden drop
- **Hard Merge – High Communication Costs – No Linear Speedup**
- But, there **is** gain
- Data still needs to fit into a single machine!

References

- *Primitives for the Manipulation of General Subdivisions and the Computation of Voronoi Diagrams* – **Guibas, L.** and **Stolfi, J.**
- *Closest-Point Problems* – **Shamos, M.I.** and **Hoey, D.**
- *On computing Voronoi diagrams by divide-prune-and-conquer* – **Amato, N.M.** and **Ramos, E.A.**
- *Chapter 10: Computational Geometry, Algorithms – Sequential and Parallel* – **Miller, R.** and **Boxer, L.**

Thank You

adarshpr@buffalo.edu

Binaries, scripts, code and results available at: <https://github.com/adrsh18/parallel>

Thanks to Dr. M Jones and CCR @ UB

Backup: Analysis

- Sequential runtime: **$O(n * \log n)$** [$T(n) = 2 * T(n/2) + O(n)$]
- “Heavy” merge step with $O(n)$. Parallelization possible?!!
- Analysis with **p** processors:
 - Each processor locally and simultaneously computes DT on $\frac{n}{p}$ points $\rightarrow O\{\frac{n}{p} * \log(\frac{n}{p})\}$
 - DTs from each processor is stitched together (happens $\log p$ times) $\rightarrow O(n * \log p)$
 - So, total runtime = $O\{\frac{n}{p} * \log(\frac{n}{p}) + n * \log p\}$
- If $p = \log n$, runtime = **$O(n \log(\log n))$**